

AI 104: Support & Operations Professionals

Building AI Literacy and Embracing the New Professional Paradigm

The Paradigm Shift



The Software Era

Humans write every rule.

Ops Example: Static FAQ pages and standard operating procedures.



The AI Era

**Humans provide examples;
machines learn the rules.**

Ops Example: Dynamic assistants that understand customer frustration and adapt.

Why AI is No Longer Optional in Operations



Scale

Human teams cannot keep up with the 24/7, global demand of modern digital business. AI acts as the first responder.



Data Overload

We generate more logs, tickets, and emails than any human can read. We need AI to triage this massive influx.



Career Resilience

Understanding AI elevates you from a task-doer to an orchestrator, significantly increasing your professional value.

AI is a field of computer science that builds systems capable of performing tasks that typically require human intelligence.

✓ Learning ✓ Reasoning ✓ Self-correction

Standard Software



Follows a rigid, set formula for predictable inputs.

Artificial Intelligence



Interprets messy, real-world data like a blurry photo or a sarcastic email.

The Capability Matrix: Weak vs. Strong AI

Feature	Weak/Narrow AI	Strong/General AI
Definition	Designed and trained for a specific task.	Full cognitive abilities of a human.
Awareness	No consciousness or genuine understanding.	Self-aware; can reason, plan, and feel.
Examples	Siri, Spam Filters, ChatGPT, Support Bots.	JARVIS, HAL 9000.
Current Status	Everywhere. This is what we use today.	Theoretical. It does not exist yet.

Directing the technology before it directs you

KEY MINDSET

AI is not replacing support and ops professionals — it's eliminating the boring parts of the job so humans can focus on the complex, creative, and human elements that machines genuinely cannot replicate.

✓ What AI Does Well

- ✓ Process large volumes of routine requests without fatigue
- ✓ Surface patterns humans would miss across thousands of data points
- ✓ Draft consistent first responses at scale
- ✓ Classify and route with high accuracy when trained correctly
- ✓ Summarize long threads and incident logs instantly

⚠ What AI Still Can't Do

- Exercise genuine empathy in complex emotional situations
- Make judgment calls that require full business context
- Handle novel edge cases it hasn't been trained on reliably
- Take accountability for outcomes
- Replace the human relationship in high-stakes escalations

Two Modes of Operational AI



Predictive AI (The Forecaster)

- Uses statistical algorithms to find if-this-then-that correlations in massive datasets.
- Doesn't create; categorizes or calculates probability.

Ops Example:

Analyzes ticket sentiment to predict if a customer is at High Risk of canceling based on past frustration.



Generative AI (The Creator)

- Learns the distribution of data from billions of examples.
- Predicts the next likely piece of data to build original output.

Ops Example:

A pro provides bullet points, and AI generates a full, formatted Standard Operating Procedure (SOP) in seconds.

Generative AI understands context, but requires explicit instruction

System Types


1 Large Language Models (Gen AI)

Understand/generate text, adapt to novel situations (but make confident mistakes).

2 Classifiers (Classic AI)

Route tickets based on historical rules. High precision on clean data.

The Context Window Funnel



The Context Window Funnel

Generative AI understands context, but requires explicit instruction

Four Core Concepts

Prompt



The instruction you give.
Quality determines output.

Metadata

Context Window



How much information the AI can hold in mind at once.

Metadata

Hallucination



Plausible-sounding but factually incorrect generated information.

Metadata

Temperature / Confidence




Controls creativity vs. consistency
(Lower = consistent; Higher = varied).

Functional

PRO TIP

Think of an LLM like a brilliant intern who has read everything ever written but has never actually done any of the work. They need supervision and explicit context.

AI intersects every stage of the modern support lifecycle



The Chatbot Efficacy Matrix: Knowing when to automate


 Ideal Conditions: When Chatbots Work Well	 Failure Modes: When Chatbots Fail
- High-volume, repetitive queries (password resets, order status)	- Complex, multi-issue problems requiring judgment
- 24/7 coverage outside business hours	- Emotionally distressed or frustrated customers
- First-level triage before human handoff	- Novel situations not in training data
- Guided troubleshooting with clear decision trees	- Nuanced billing or legal disputes

Catching operational issues earlier and resolving them faster

OPS LIFECYCLE



REAL SCENARIO



AI Manifested: The Operations Toolkit



Patterns in Data

Concept: Finding the needle in the haystack.

Ops Use: Cybersecurity login anomaly detection.



Natural Language Processing (NLP)

Concept: Understanding human-speak.

Ops Use: Auto-tagging and routing support tickets.



Internet of Things (IoT)

Concept: Connecting physical objects to the internet.

Ops Use: Smart Factories utilizing predictive maintenance.



Generative Systems

Concept: Foundational Large Language Models.

Ops Use: Creating Synthetic Data for safe system testing.

AI 104: Support & Operations Professionals

Artificial Intelligence and
Machine Learning



AI models are mathematical algorithms that learn patterns from data to make decisions—without being explicitly programmed.



Takes Input Data

Raw information is fed into the system.



Learns Patterns

The algorithm identifies hidden structures during the training phase.



Produces Output

The system generates a prediction, decision, or generated result.

Three Core Learning Styles

Supervised Learning



Data Type:
Labeled Data (Input
+ Correct Output)

Goal: Predict correct answers.

Real-World Example: Email
spam detection (Gmail).

Unsupervised Learning



Data Type:
Unlabeled Data

Goal: Find hidden patterns or
groups.

Real-World Example: Customer
segmentation based on buying
behavior (Amazon).

Reinforcement Learning




Data Type:
Trial & Error
(Rewards/Penalties)

Goal: Maximize reward over
time.

Real-World Example: Game-
playing AI learning winning
strategies (AlphaGo).

Primary Task Categories

Classification Models



Output:
Categories / Labels

Example: Fraud detection in PayPal
(Transaction → Fraud / Not Fraud)

Regression Models



Output: Continuous values

Example: House price prediction on Zillow
(Predicts exact price based on features)

Clustering Models



Output: Groups of similar data

Example: Music grouping in Spotify
(Clusters songs based on listening patterns)

Recommendation Models



Output: Personalized suggestions based on behavior

Example: Movie suggestions on Netflix

Algorithm Gallery: Classification Models

Classification

Decision Tree

Classify using a series of if-then rules. Easy to explain to non-technical stakeholders.

Key words: rules-based, explainable, medical diagnosis, loan approval

Classification

Random Forest

Many decision trees voted together. More accurate than a single tree, less prone to overfitting.

Key words: high accuracy, tabular data, fraud detection, customer churn

Classification

KNN (K-Nearest Neighbours)

Classify by looking at the K most similar past examples. Simple, intuitive, no training phase.

Key words: similarity-based, recommend similar items, small dataset

Classification

Logistic Regression

Classify into 2 categories (yes/no). Despite its name, it's a classifier.

Key words: spam or not, churn or not, binary outcome

Algorithm Gallery: Advanced Models

Regression

Linear Regression

Predict a continuous number when you expect a straight-line relationship.

Key words: predict price, forecast sales, estimate score, how much/many

Clustering

K-Means Clustering

Group unlabelled data into K natural groups.

Key words: group customers, segment users, find patterns, no labels

Neural Network

Neural Network / Deep Learning

Complex patterns in images, text, audio. Needs lot often used for dad compute.

Key words: image recognition, speech, text generation, large datasets


Regression

Polynomial / Multiple Regression

Predict a number using multiple input features, or when the relationship curves (non-linear).

Key words: multiple factors, curved trend, house price with many features

The Complexity Evolution




The AI Taxonomy: Master Cheat Sheet

Classification Type	Core Focus	Real-World Example
Supervised	Labeled data	Spam detection
Unsupervised	Pattern discovery	Customer groups
Reinforcement	Reward-based learning	Game AI
Classification	Categories	Fraud detection
Regression	Numbers	Price prediction
Deep Learning	Complex patterns	Face recognition
Generative AI	Content creation	Chatbots

Machine Learning: The Engine Inside the AI

Artificial Neural Networks (ANN) mimic human neurons, training themselves without fixed instructions.




The goal of an
AI model is not
to memorize
the past...



**...but to learn
the patterns
necessary
to predict
the future.**

The ultimate challenge of model training is balancing how
well a system learns without simply memorizing data.




Underfitting (The Simple Failure)

The Flaw: Model is too simple;
fails to learn the pattern.

Performance: Fails on Training
Data ❌ | Fails on New Data ❌

Root Causes: Not enough training
time, missing critical features.



Overfitting (The Complex Failure)

The Flaw: Model learns too much,
including pure noise.

Performance: Perfect on Training
Data ✅ | Fails on New Data ❌

Root Causes: Model is too
complex, dataset is too small.

Real-World Impact: Predicting House Prices



Underfit

Only looks at the number of rooms. Ignores location, size, and amenities. Predictions fail everywhere.



Overfit


Obsesses over random noise, like a specific crack in the driveway of one training house. Cannot predict prices in a new neighborhood.




Perfect Fit

Identifies the true underlying patterns: square footage, zip code, and market trends.


Part 1: The Neuronal Core



Part 2: The Correction Loop



Part 3: The Hyperparameters



Learning Rate (The Step Size)

Controls how drastically weights are updated.

Too high: Overshoots the solution.

Too low: Training takes forever.



Epoch (The Repetitions)

One complete pass through the entire dataset.

Multiple passes improve learning, but too many trigger **Overfitting**.




Batch Size (The Chunks)

Number of samples processed at once.


Small batch: Faster but chaotic updates.

Large batch: Stable but slower learning.

Synthesis: The Model Training Lifecycle



Filtering the Noise: Activation Bias



The Manager Analogy


Setting a Filter:

Don't bother me with minor complaints; only come to me if the total loss is over \$500.

The Application:

The \$500 is the Bias. It filters out the noise so the AI system stays focused entirely on what matters.

The Probability Engine: How LLMs Actually Work



Key Takeaway: LLMs operate through pattern recognition and probability rather than factual reasoning or conscious thought.

AI 104: Support & Operations Professionals

Prompt Engineering

Prompt Engineering: The Steering Wheel

Context: Operating in a world where technology generates new data (text, code, images).



Few-Shot Learning in Action (Claude AI)

Scenario: Standardizing messy incident reports.

The Prompt: Turn raw notes into formal reports. Follow this style exactly:

Raw Data'

Server went down at 2pm...

'Formal Output'

[INCIDENT-001] | Status: Resolved | Duration: 60m...

Beyond Text: Visual and Voice Operations



Text-to-Image (T2I) Prototyping

The Shift: Moving from hours searching stock art or waiting for designers to creating UI mockups in seconds.

Application: Generating instant visual prototypes, like Figma AI layouts or DALL-E server rack visualizations, directly from plain English descriptions.




Multimodal Voice (The Eyes-On Assistant)

Scenario: Infrastructure engineer with hands occupied at a server rack.

User: Status light on primary PSU is flashing green, secondary is solid amber on an R740. Hot-swap or firmware sync?

AI: Solid amber indicates predictive failure or power mismatch. Prep a replacement unit for a hot-swap immediately.

The anatomy of a production-grade AI prompt



✗ Weak Prompt Example

Write a reply to this angry customer.

✓ Strong Prompt Example

```
[ROLE]: You are a senior support agent at a cloud SaaS company.  
[TASK]: Draft a response to the following ticket...  
[CONTEXT]: ...from a customer who has been waiting 3 days for a refund.  
[CONSTRAINTS]: Tone: empathetic and apologetic. Acknowledge delay, explain timeline (5-7 days). Do not promise anything beyond policy.  
[FORMAT]: Keep it under 150 words.
```

Diving into ChatGPT-5

ChatGPT-5 is your **"Logic Engine."** It is the most advanced model for multi-step reasoning and complex problem-solving.



Key Feature

Deep Thinking Mode

Allows the model to pause and "reason" through a problem before answering.



Ops Use Case

Perfect for **complex debugging** (e.g., microservices logs) or writing **sophisticated scripts**.



Power User Tip

272k Token Window

Feed it an entire codebase or a 200-page technical manual in one go.

Diving into Chat GPT-5

Use ChatGPT-5 as your "Senior Architect." It excels at complex reasoning and connecting disparate pieces of information.

Example A: Script Debugging

- **Prompt:** "You are a Senior Systems Engineer. I have a Python script that is supposed to automate user offboarding in Azure, but it's failing at the 'Revoke License' step with a 403 error. Here is the code: [Paste Code]. Analyze the logic, identify the most likely permission mismatch, and provide the corrected snippet."

GPT-5 Response (Simulated Output):

Thinking (GPT-5 Reasoning Engine)

- User is performing Azure AD (Entra ID) offboarding.
- Error is 403 Forbidden during the license revocation step.
- Script uses MS Graph API via the msal library.
- **Common Pitfalls:** 403 usually implies the Service Principal lacks Directory.ReadWrite.All or User.ReadWrite.All. However, for *licenses*, specifically, Organization.Read.All is often missed.
- **Code Check:** The script calls `/users/{id}/assignLicense`. This requires User.ReadWrite.All **AND** Directory.ReadWrite.All.
- **Logic Check:** If the user is a Global Admin, a standard Service Principal cannot offboard them without RoleManagement.ReadWrite.Directory.

Diving into Chat GPT-5

Analysis: The 403 Forbidden Error

Based on the code provided, the 403 error isn't a syntax bug; it's a **Scope Mismatch**.

The Root Cause: Your Service Principal is likely missing the **Organization.Read.All** permission. While User.ReadWrite.All lets you edit the user, the license assignment endpoint needs to "read" the organization's SKU details to verify which licenses are being removed.

Required API Permissions (Microsoft Graph):

1. User.ReadWrite.All (Application)
2. Directory.ReadWrite.All (Application)
3. **Organization.Read.All** (Missing piece)

Diving into Chat GPT-5

Corrected Code Snippet (Logic Update)

Ensure your header includes the correct scope and check for administrative locks before the call:

```
# Ensure the Service Principal has 'Organization.Read.All'
def revoke_licenses(user_id, access_token):
    endpoint = f"https://graph.microsoft.com/v1.0/users/{user_id}/assignLicense"
    headers = {'Authorization': f'Bearer {access_token}', 'Content-Type': 'application/json'}

    # Correct Payload for revoking all licenses
    # We must fetch active licenses first to populate 'removedLicenses'
    payload = {
        "addLicenses": [],
        "removedLicenses": active_sku_ids # Ensure this list is populated
    }

    response = requests.post(endpoint, headers=headers, json=payload)
    if response.status_code == 403:
        print("Error: 403. Check 'Organization.Read.All' permissions in Azure Portal.")
    return response.json()
```

Diving into Chat GPT-5

Example B: Root Cause Analysis (RCA)

Prompt: "I am drafting an RCA for a 2-hour outage in our Slack-to-Jira integration. Based on these raw incident logs: [Paste Logs], write a professional post-mortem. Structure it with: Executive Summary, Technical Root Cause, and 3 Preventative Actions. Use a tone that is accountable yet confident."

GPT-5 Response (Simulated Output):

Incident Report: Slack-Jira Integration Outage

Date: April 21, 2026

Duration: 120 Minutes (14:00 - 16:00 UTC)

Status: Resolved

1. Executive Summary

On April 21, 2026, the bi-directional sync between Slack and Jira experienced a total service disruption for approximately two hours. This prevented automated ticket creation from Slack channels and disrupted real-time status updates. Service was restored by rolling back a faulty webhook listener configuration. No data was lost, though a backlog of 450+ events required manual reconciliation.

Diving into Chat GPT-5

Example B: Root Cause Analysis (RCA)

Prompt: "I am drafting an RCA for a 2-hour outage in our Slack-to-Jira integration. Based on these raw incident logs: [Paste Logs], write a professional post-mortem. Structure it with: Executive Summary, Technical Root Cause, and 3 Preventative Actions. Use a tone that is accountable yet confident."

2. Technical Root Cause

Analysis of the logs (specifically Error Code: E_TIMED_OUT at 14:04:12) indicates a **Webhook Deadlock**.

- **Trigger:** A minor update to the Amazon Transform middleware introduced a recursive loop where a Jira status update triggered a Slack notification, which the middleware incorrectly re-interpreted as a new Jira update.
- **Result:** This created an exponential surge in API calls, hitting Slack's rate limits (429 Too Many Requests) and ultimately crashing the listener pod due to memory exhaustion (OOMKill).

3. Preventative Actions

- **Idempotency Keys:** Implement mandatory idempotency headers for all Amazon Transform traffic to ignore duplicate event IDs.
- **Circuit Breaker Implementation:** Deploy a circuit breaker pattern that automatically kills the webhook sync if the error rate exceeds 15% in a 60-second window.
- **Enhanced Monitoring:** Configure **Amazon Q** to alert on "Recursive Traffic Patterns" rather than just standard uptime pings.

Exploring Microsoft Copilot

Copilot is the "**Corporate Navigator.**" It excels at finding information buried in internal Microsoft Graph (Teams, Outlook, SharePoint).



Ops Use Case

Internal Search

"Find the SOP for 'Firewall Exception Requests' that was shared in the 'Security-Ops' Teams channel last month."



Feature

Copilot Agents

Create mini-AI assistants in Teams to handle specific tasks, like onboarding a new hire or checking ticket status.



Integration

Microsoft 365

Seamless connectivity across Outlook, Teams, and SharePoint to unify your organizational knowledge.

AI 104: Support & Operations Professionals


Ethics and Governance in
Artificial Intelligence

Governance workflow: Maintaining 100% of human accountability

THE GOLDEN RULE: AI drafts. Humans approve.

(Captures **80%** of **efficiency** gains while maintaining **100% of accountability**).
Over-relying on AI responses without human review is the most common mistake.

Governance workflow: Maintaining 100% of human accountability



✓ Review Checklist (What to look for)

- ✓ **Factual accuracy:** Right product name, price, policy?
- ✓ **Tone calibration:** Appropriate for this specific situation?
- ✓ **Completeness:** Did it address all parts of the question?
- ✓ **Hallucinations:** Did it invent features or commitments?
- ✓ **Personalization:** Does it read like a robot?


⚠ When NOT to use AI

- ⚠ Legal/compliance issues, high-stakes emotional escalations (loss/grief/anger), executive communications, or when AI lacks the required data.

The Ethical Imperative: Tool vs. Behavior

The New Stakes

- 1. **The Black Box:** Creators cannot always explain the output.
- 2. **Massive Scale:** Instant generation of text, images, and deepfakes.



The Tool: Technology is entirely neutral until it is applied. We cannot simply blame the algorithm.

The Behavior: How we deploy it—the guardrails we remove and the data we feed it—**defines** the ethical outcome.

The Mandate: Move from asking “Can we build this?” to “Should we build this?”

The Ethical AI Framework

Use these three pillars to evaluate the integrity of any AI initiative.

01 Intent

(The Why)



Is the primary goal to help the user, or is it to manipulate behavior?

02 Process

(The How)



How was the model built? Does it have Data Hygiene (removing sensitive identifiers) and a paper trail for Auditability?

03 Outcome

(The What)



Who wins and who loses? Does it create a digital divide? Are users notified they are interacting with AI (Transparency)?

The real-world tools landscape across Support and Ops

Vendor Matrix			
	Category	Platform	Best For
1	Helpdesk (Support)	Zendesk AI, Freshdesk Freddy AI	Core support teams and ticket routing
2	CRM (Both)	Salesforce Einstein	Enterprise CRM + service cloud
3	Incident Mgmt (Ops)	PagerDuty AIOps	On-call teams, SRE routing
4	Observability (Ops)	Datadog Watchdog, Dynatrace Davis AI	Infrastructure APM, Enterprise root cause
5	Conversational AI (Support)	Intercom Fin	Automated customer resolutions
6	Knowledge Mgmt (Both)	Notion AI, Confluence AI	KB authoring and internal search
7	General AI (Both)	Microsoft Copilot, Claude/ChatGPT	Drafting, documentation, analysis

Approved AI Directory — Core Models & Cloud Platforms

Authorized foundational large language models (LLMs) and enterprise infrastructure platforms for general knowledge work and secure deployment.

Conversational AI & Workspaces

General Chat & Reasoning



Gemini for Google Workspace

Company-wide license; everyone can access.



ChatGPT-5

Advanced reasoning and drafting capabilities.



Claude AI

Large-context analysis and document summarization.



Enterprise Cloud Ecosystems

Infrastructure & API Access



AzureOpenAI

Secure enterprise deployment and integration.



Vertex AI

Custom model building and machine learning operations.



Amazon Bedrock

Scalable foundation model API access.



Approved AI Directory — Specialized & Collaborative Tools

Authorized point-solutions for engineering, design, and team productivity workflows.

Engineering & Development

Code & Tech Ops



GitHub Copilot

AI pair programmer.



Cursor

AI-first code editor.



Amazon Q & Amazon Transform

Business/Developer assistant and code modernization.



Design & Team Collaboration

Assets & Knowledge



Figma AI

Design generation and asset management.



Rovo - Atlassian Intelligence

Enterprise search and knowledge graphing.



Slack

Thread summarization and workflow automation.



Productivity & Utility

Translation & Whiteboarding



DeepL

High-accuracy neural machine translation.



Miro



Approved with condition. Users **MUST** configure settings using the "opt-out" option to prevent corporate data from being used for AI training.

Choosing the Right Tool

1. Where does my data go (training models)? 

2. Can I see why the AI made a decision (explainability)? 

3. How do I correct the AI when it's wrong? 

4. What does the human override look like? 

Guardrails, ethics, and professional responsibility

Risk Domain

Hallucinations

AI fabricates terms. Always verify authoritative sources.

Risk Domain

Human-in-the-Loop

Refunds, terminations, legal disputes require judgment.

Risk Domain

Bias

Inherits historical biases (e.g., ticket routing favoring demographics). Audit disparate impact.

Risk Domain

Compliance

GDPR, CCPA, HIPAA apply to AI-processed data.

Risk Domain

Data Privacy

Never paste PII or credentials into public tools.

Risk Domain

Auditability

Can you explain the decision? Maintain human decision audit trails.

The Practical Checklist

- ☒ Know your organization's approved AI tools.
- ☒ Never input SSNs, card numbers, or full records without IT approval.
- ☒ Treat every AI output as a draft until verified.
- ☒ Report AI errors to improve the models.
- ☒ Know your escalation path for scenarios bypassing AI entirely.

AI 104: Support & Operations Professionals

What's Next ?

The Capability Blueprint: From Reactive to Autonomous



RAG (Retrieval-Augmented Generation)

Look up info

Shift: From guessing to fetching real-time, external data.

Example: Searching company leave policies accurately.



Agentic AI

Act independently

Shift: From AI that talks (Chatbot) to AI that does work.

Example: Reading emails and booking meetings automatically.



Agentic RAG

Think, Search & Act

Shift: Multi-step reasoning combining iterative searches with decision-making.

Example: An AI research assistant combining insights from multiple sources.

Shaping the 2025–2026 Enterprise Landscape



Agentic Workflows

End-to-end autonomous execution (e.g., coding to deployment). Expected in 40% of enterprise apps by 2026.



Small Language Models (SLMs)

Domain-specific, faster, and highly cost-efficient.



Multimodal AI

Seamlessly processing text, image, audio, and video together.



Human + AI Collaboration

Evolving from a simple tool into an active digital co-worker.



Hyperautomation

Autonomous execution of entire business processes (e.g., supply chain decisions).



AI Governance & Safety

Expanding focus on risk management, compliance, and ethics.



AI-native Systems

Software architected with AI at its core from the design stage.



AI is evolving from answering questions to doing complete tasks intelligently.



AI 104: Support & Operations Professionals


Mastering Daily Integration

Driving adoption: From execution to oversight

The Cultural Shift: Will AI take my job?

Roles are shifting from answering tickets to managing AI that answers tickets.
The goal is **AI-augmented professionals**, not **AI-dependent** ones.

Driving adoption: From execution to oversight



The Skill Value Shift Matrix

↓ Declining in Value

~~Manual ticket sorting / copy-paste triage~~
~~Basic knowledge article retrieval~~
~~Template-based first responses~~
~~Manual log parsing / grep debugging~~
~~Static threshold-based alerts~~

↑ Rising in Value

- ✓ **Prompt engineering & AI workflow design**
- ✓ **AI output quality assurance**
- ✓ **Process design for human-AI collaboration**
- ✓ **AI governance and ethics judgment**
- ✓ **Interpreting AI-generated analytics**

Axiom 01

AI is a tool, not a replacement.

AI's Domain	Human's Domain
Handles Volume. Drafting, pattern matching, baseline generation.	Handles Nuance. Judgment, empathy, contextual understanding.



Axiom 02

Garbage in, garbage out.

$$\left[\begin{array}{c} \mathbf{60} \\ \text{Seconds} \end{array} \right] = \left[\begin{array}{c} \mathbf{10} \\ \text{Minutes} \end{array} \right]$$

Invested in a better prompt.

Saved in editing.

The quality of your prompt determines the quality of your output.

Axiom 03


Always verify before sending.



The Strategic Mindset: Navigating Long-Term Impact

Axiom 04


Understand AI's failure modes.



Knowing these risks makes you a safer and more effective AI user—not a paranoid one.

Axiom 05

Learn AI to outperform.



Start building the skill now. The professionals who learn AI will outperform those who don't.

Link to L&D Catalog:-

<https://sites.google.com/globallogic.com/ld-global-offering/>

Thank you